

Workshop5 Output:

```
✓ 15s [1] import os
import random
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Rescaling
from tensorflow.keras.preprocessing.image import load_img, img_to_array
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
from sklearn.metrics import classification_report
```

```
✓ 1m [4] from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
✓ 0s [2] train_dir = "/content/drive/MyDrive/Workshop5_AI/FruitinAmazon/train"
test_dir = "/content/drive/MyDrive/Workshop5_AI/FruitinAmazon/test"
```

```
✓ 0s [5] class_names = os.listdir(train_dir)
print(f"Classes: {class_names}")
```

Classes: ['guarana', 'graviola', 'acai', 'cupuacu', 'tucuma', 'pupunha']

```
▶ def visualize_images(train_dir, class_names):
    fig, axes = plt.subplots(2, len(class_names) // 2, figsize=(12, 6))
    axes = axes.flatten()
    for i, class_name in enumerate(class_names):
        class_path = os.path.join(train_dir, class_name)
        img_name = random.choice(os.listdir(class_path))
        img_path = os.path.join(class_path, img_name)
        img = load_img(img_path)
        axes[i].imshow(img)
        axes[i].set_title(class_name)
        axes[i].axis("off")
    plt.show()

visualize_images(train_dir, class_names)
```

guarana



graviola



acai



cupuacu



tucuma




pupunha



```
[7] damagedImages = []
    for class_name in class_names:
        class_path = os.path.join(train_dir, class_name)
        for img_name in os.listdir(class_path):
            img_path = os.path.join(class_path, img_name)
            try:
                img = load_img(img_path) # Try opening the image
            except (IOError, SyntaxError):
                damagedImages.append(img_path)
                os.remove(img_path)
                print(f"Damaged image removed: {img_path}")

    if not damagedImages:
        print("No Damaged Images Found.")
```

 No Damaged Images Found.

```
[8] img_height, img_width = 128, 128
    batch_size = 32
    validation_split = 0.2
```

```
[9] train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    labels='inferred',
    label_mode='int',
    image_size=(img_height, img_width),
    batch_size=batch_size,
    shuffle=True,
    validation_split=validation_split,
    subset='training',
    seed=123
)
```

➞ Found 90 files belonging to 6 classes.
Using 72 files for training.

```
[10] val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    labels='inferred',
    label_mode='int',
    image_size=(img_height, img_width),
    batch_size=batch_size,
    shuffle=False,
    validation_split=validation_split,
    subset='validation',
    seed=123
)
```

➞ Found 90 files belonging to 6 classes.
Using 18 files for validation.

```
[11] rescale = Rescaling(1./255)
    train_ds = train_ds.map(lambda x, y: (rescale(x), y))
    val_ds = val_ds.map(lambda x, y: (rescale(x), y))
```

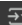
```
[12] num_classes = len(class_names)

model = Sequential([
    Conv2D(32, (3,3), activation='relu', padding='same', input_shape=(img_height, img_width, 3)),
    MaxPooling2D((2,2), strides=2),

    Conv2D(32, (3,3), activation='relu', padding='same'),
    MaxPooling2D((2,2), strides=2),

    Flatten(),
    Dense(64, activation='relu'),
    Dense(128, activation='relu'),
    Dense(num_classes, activation='softmax')
])

model.summary()
```

 /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/'input_dim' super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 32)	996
max_pooling2d (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_1 (Conv2D)	(None, 64, 64, 32)	9,248
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 32)	0
flatten (Flatten)	(None, 32768)	0
dense (Dense)	(None, 64)	2,100,224
dense_1 (Dense)	(None, 128)	8,320
dense_2 (Dense)	(None, 6)	774

Total params: 2,116,454 (8.07 MB)
Trainable params: 2,116,454 (8.07 MB)
Non-trainable params: 0 (0.00 B)

```
[13] model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```

```
[14] callbacks = [
    ModelCheckpoint("best_model.h5", save_best_only=True, monitor="val_accuracy", mode="max"),
    EarlyStopping(monitor="val_loss", patience=10, restore_best_weights=True)
]

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=250,
    batch_size=16,
    callbacks=callbacks
)
```

```
Epoch 1/250
3/3 0s 375ms/step - accuracy: 0.1904 - loss: 2.2249WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save_model()'
Epoch 2/250
3/3 4s 694ms/step - accuracy: 0.1879 - loss: 2.2723 - val_accuracy: 0.8333 - val_loss: 1.1333
Epoch 3/250
3/3 3s 813ms/step - accuracy: 0.1853 - loss: 1.8386 - val_accuracy: 0.0000e+00 - val_loss: 1.7819
Epoch 4/250
3/3 4s 536ms/step - accuracy: 0.3338 - loss: 1.6873 - val_accuracy: 0.0000e+00 - val_loss: 1.8785
Epoch 5/250
3/3 2s 475ms/step - accuracy: 0.1793 - loss: 1.6343 - val_accuracy: 0.0000e+00 - val_loss: 1.9184
Epoch 6/250
3/3 3s 467ms/step - accuracy: 0.3702 - loss: 1.5101 - val_accuracy: 0.3333 - val_loss: 1.7069
Epoch 7/250
3/3 2s 552ms/step - accuracy: 0.5425 - loss: 1.4167 - val_accuracy: 0.1667 - val_loss: 1.7707
Epoch 8/250
3/3 3s 709ms/step - accuracy: 0.5000 - loss: 1.2407 - val_accuracy: 0.4444 - val_loss: 1.5451
Epoch 9/250
3/3 4s 530ms/step - accuracy: 0.6302 - loss: 1.0921 - val_accuracy: 0.5000 - val_loss: 1.2680
Epoch 10/250
3/3 2s 462ms/step - accuracy: 0.6819 - loss: 0.8641 - val_accuracy: 0.8333 - val_loss: 0.8749
Epoch 11/250
3/3 2s 521ms/step - accuracy: 0.7934 - loss: 0.6959 - val_accuracy: 0.2222 - val_loss: 1.5155
Epoch 12/250
3/3 3s 566ms/step - accuracy: 0.7756 - loss: 0.6744 - val_accuracy: 0.5000 - val_loss: 1.3197
Epoch 13/250
3/3 0s 604ms/step - accuracy: 0.8443 - loss: 0.6197WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save_model()'
Epoch 14/250
3/3 4s 842ms/step - accuracy: 0.8485 - loss: 0.6031 - val_accuracy: 0.8889 - val_loss: 0.8000
Epoch 15/250
3/3 2s 545ms/step - accuracy: 0.9301 - loss: 0.4003 - val_accuracy: 0.8333 - val_loss: 0.9476
Epoch 16/250
3/3 2s 546ms/step - accuracy: 0.9332 - loss: 0.3115 - val_accuracy: 0.7778 - val_loss: 0.7349
Epoch 17/250
3/3 2s 475ms/step - accuracy: 0.9813 - loss: 0.1743 - val_accuracy: 0.8333 - val_loss: 0.6318
Epoch 18/250
3/3 2s 690ms/step - accuracy: 1.0000 - loss: 0.1146 - val_accuracy: 0.8333 - val_loss: 0.5765
Epoch 19/250
3/3 3s 778ms/step - accuracy: 0.9891 - loss: 0.0582 - val_accuracy: 0.7778 - val_loss: 0.8537
Epoch 20/250
3/3 2s 536ms/step - accuracy: 1.0000 - loss: 0.0506 - val_accuracy: 0.8889 - val_loss: 0.5510
Epoch 21/250
3/3 0s 358ms/step - accuracy: 1.0000 - loss: 0.0226WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save_model()'
Epoch 22/250
3/3 2s 575ms/step - accuracy: 1.0000 - loss: 0.0232 - val_accuracy: 0.9444 - val_loss: 0.4446
Epoch 23/250
3/3 2s 547ms/step - accuracy: 1.0000 - loss: 0.0240 - val_accuracy: 0.8889 - val_loss: 0.3478
Epoch 24/250
3/3 2s 474ms/step - accuracy: 1.0000 - loss: 0.0142 - val_accuracy: 0.9444 - val_loss: 0.4080
Epoch 25/250
3/3 2s 514ms/step - accuracy: 1.0000 - loss: 0.0101 - val_accuracy: 0.8889 - val_loss: 0.7530
Epoch 26/250
3/3 4s 789ms/step - accuracy: 0.9627 - loss: 0.0584 - val_accuracy: 0.8889 - val_loss: 0.4801
Epoch 27/250
3/3 2s 538ms/step - accuracy: 1.0000 - loss: 0.0071 - val_accuracy: 0.7778 - val_loss: 0.5452
Epoch 28/250
3/3 3s 533ms/step - accuracy: 1.0000 - loss: 0.0118 - val_accuracy: 0.8333 - val_loss: 0.5479
Epoch 29/250
3/3 2s 476ms/step - accuracy: 1.0000 - loss: 0.0079 - val_accuracy: 0.8889 - val_loss: 0.5369
Epoch 30/250
3/3 2s 537ms/step - accuracy: 1.0000 - loss: 0.0032 - val_accuracy: 0.8889 - val_loss: 0.5783
Epoch 31/250
3/3 2s 537ms/step - accuracy: 1.0000 - loss: 0.0037 - val_accuracy: 0.8889 - val_loss: 0.6067
Epoch 32/250
3/3 4s 766ms/step - accuracy: 1.0000 - loss: 0.0034 - val_accuracy: 0.9444 - val_loss: 0.6134
Epoch 33/250
3/3 2s 477ms/step - accuracy: 1.0000 - loss: 0.0026 - val_accuracy: 0.9444 - val_loss: 0.6048

test_ds = tf.keras.preprocessing.image_dataset_from_directory(
    test_dir,
    labels='inferred',
    label_mode='int',
    image_size=(img_height, img_width),
    batch_size=batch_size,
    shuffle=False
)

test_ds = test_ds.map(lambda x, y: (rescale(x), y))
test_loss, test_accuracy = model.evaluate(test_ds)
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")

Found 30 files belonging to 6 classes.
1/1 10s 10s/step - accuracy: 0.5667 - loss: 1.3295
Test Accuracy: 56.67%

[16] model.save("final_model.h5")
loaded_model = tf.keras.models.load_model("final_model.h5")

WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save_model(model)'. This file format is considered legacy. We recommend using instead
WARNING:absl:compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile_metrics' will be empty until you train or evaluate the model.
```

```

y_true = []
y_pred = []

for images, labels in test_ds:
    preds = loaded_model.predict(images)
    y_pred.extend(np.argmax(preds, axis=1))
    y_true.extend(labels.numpy())

print(classification_report(y_true, y_pred, target_names=class_names))

```

1/1 ————— 0s 343ms/step

	precision	recall	f1-score	support
guarana	0.75	0.60	0.67	5
graviola	0.33	0.40	0.36	5
acai	0.67	0.80	0.73	5
cupuacu	0.75	0.60	0.67	5
tucuma	0.71	1.00	0.83	5
pupunha	0.00	0.00	0.00	5
accuracy			0.57	30
macro avg	0.54	0.57	0.54	30
weighted avg	0.54	0.57	0.54	30

```

plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

```

