

## Worksheet 6 output:

```
46 import os
import random
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, BatchNormalization
from tensorflow.keras.preprocessing.image import load_img, ImageDataGenerator
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.regularizers import l2
from sklearn.metrics import classification_report
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.layers import GlobalAveragePooling2D, Input
from tensorflow.keras.models import Model

[4] from google.colab import drive
drive.mount('/content/drive')

↗ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Task-1

train_dir = "/content/drive/MyDrive/Workshop5_AI/FruitInAmazon/train"
test_dir = "/content/drive/MyDrive/Workshop5_AI/FruitInAmazon/test"

[6] class_names = os.listdir(train_dir)
print(f"Classes: {class_names}")

↗ Classes: ['guarana', 'graviola', 'acai', 'cupuacu', 'tucuma', 'pupunha']

def visualize_images(train_dir, class_names):
    fig, axes = plt.subplots(2, len(class_names) // 2, figsize=(12, 6))
    axes = axes.flatten()
    for i, class_name in enumerate(class_names):
        class_path = os.path.join(train_dir, class_name)
        img_name = random.choice(os.listdir(class_path))
        img_path = os.path.join(class_path, img_name)
        img = load_img(img_path)
        axes[i].imshow(img)
        axes[i].set_title(class_name)
        axes[i].axis("off")
    plt.show()

visualize_images(train_dir, class_names)
```

guarana



graviola



acai



cupuacu



tucuma



pupunha



```
▶ damagedImages = []
for class_name in class_names:
    class_path = os.path.join(train_dir, class_name)
    for img_name in os.listdir(class_path):
        img_path = os.path.join(class_path, img_name)
        try:
            img = load_img(img_path) # Try opening the image
        except (IOError, SyntaxError):
            damagedImages.append(img_path)
            os.remove(img_path)
            print(f"Damaged image removed: {img_path}")

if not damagedImages:
    print("No Damaged Images Found.")
```

📁 No Damaged Images Found.

```
[9] img_height, img_width = 128, 128
    batch_size = 32
    validation_split = 0.2
```

```

▶ train_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=validation_split,
    rotation_range=30,
    width_shift_range=0.2,
    height_shift_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    brightness_range=[0.8, 1.2]
)

val_datagen = ImageDataGenerator(rescale=1./255, validation_split=validation_split)

```

```

[11] train_ds = train_datagen.flow_from_directory(
    train_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='sparse',
    subset='training',
    shuffle=True,
    seed=123
)

```

Found 72 images belonging to 6 classes.

```

0s ▶ val_ds = val_datagen.flow_from_directory(
    train_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='sparse',
    subset='validation',
    shuffle=False,
    seed=123
)

```

Found 18 images belonging to 6 classes.

```

0s [13] num_classes = len(class_names)

model = Sequential([
    Conv2D(32, (3,3), activation='relu', padding='same', kernel_regularizer=l2(0.001), input_shape=(img_height, img_width, 3)),
    BatchNormalization(),
    Conv2D(64, (3,3), activation='relu', padding='same', kernel_regularizer=l2(0.001)),
    BatchNormalization(),
    MaxPooling2D((2,2)),
    Dropout(0.25),

    Conv2D(128, (3,3), activation='relu', padding='same', kernel_regularizer=l2(0.001)),
    BatchNormalization(),
    MaxPooling2D((2,2)),
    Dropout(0.4),

    Flatten(),
    Dense(256, activation='relu', kernel_regularizer=l2(0.001)),

```


```
BatchNormalization(),
Dropout(0.5),
Dense(num_classes, activation='softmax')
})

model.summary()
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an 'input\_shape'/'input\_dim' argument to a layer. When using Sequential, use super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 16)	960
batch_normalization (BatchNormalization)	(None, 128, 128, 16)	128
conv2d_1 (Conv2D)	(None, 128, 128, 64)	16,400
batch_normalization_1 (BatchNormalization)	(None, 128, 128, 64)	256
max_pooling2d (MaxPooling2D)	(None, 64, 64, 64)	0
dropout (Dropout)	(None, 64, 64, 64)	0
conv2d_2 (Conv2D)	(None, 64, 64, 128)	73,808
batch_normalization_2 (BatchNormalization)	(None, 64, 64, 128)	512
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 128)	0

	dropout_1 (Dropout)	(None, 32, 32, 128)	0
	flatten (Flatten)	(None, 131072)	0
	dense (Dense)	(None, 256)	33,554,688
	batch_normalization_3 (BatchNormalization)	(None, 256)	1,024
	dropout_2 (Dropout)	(None, 256)	0
	dense_1 (Dense)	(None, 6)	1,542
Total params: 33,651,398 (128.37 MB)			
Trainable params: 33,650,438 (128.37 MB)			
Non-trainable params: 960 (3.75 KB)			

```
[14] model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```

```
callbacks = [
    ModelCheckpoint("best_model.h5", save_best_only=True, monitor="val_accuracy", mode="max"),
    EarlyStopping(monitor="val_loss", patience=5, restore_best_weights=True),
    ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e-6)
]

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=30,
    batch_size=16,
    callbacks=callbacks
)
```

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data\_adapters/py\_dataset\_adapter.py:121: UserWarning: Your 'PyDataset' class should call 'super().\_\_init\_\_(\*\*kwargs)'

Epoch 1/30  
3/3 — 0s 4s/step - accuracy: 0.2171 - loss: 3.5672WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This format is deprecated. Use the Keras 2.x HDF5 format instead.  
3/3 — 29s 8s/step - accuracy: 0.2288 - loss: 3.5763 - val\_accuracy: 0.2778 - val\_loss: 2.3253 - learning\_rate: 0.0010

Epoch 2/30  
3/3 — 16s 7s/step - accuracy: 0.4938 - loss: 2.3590 - val\_accuracy: 0.1667 - val\_loss: 2.7081 - learning\_rate: 0.0010

Epoch 3/30  
3/3 — 16s 4s/step - accuracy: 0.6033 - loss: 2.3251 - val\_accuracy: 0.2222 - val\_loss: 3.2056 - learning\_rate: 0.0010

Epoch 4/30  
3/3 — 0s 4s/step - accuracy: 0.5676 - loss: 2.3569WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This format is deprecated. Use the Keras 2.x HDF5 format instead.  
3/3 — 21s 7s/step - accuracy: 0.5576 - loss: 2.3870 - val\_accuracy: 0.3333 - val\_loss: 3.7672 - learning\_rate: 0.0010

Epoch 5/30  
3/3 — 37s 4s/step - accuracy: 0.4868 - loss: 2.9064 - val\_accuracy: 0.3333 - val\_loss: 4.9269 - learning\_rate: 5.0000e-04

Epoch 6/30  
3/3 — 17s 4s/step - accuracy: 0.6128 - loss: 2.6247 - val\_accuracy: 0.2778 - val\_loss: 6.4192 - learning\_rate: 5.0000e-04

```

176 test_datagen = ImageDataGenerator(rescale=1./255)
177 test_ds = test_datagen.flow_from_directory(
178     test_dir,
179     target_size=(img_height, img_width),
180     batch_size=batch_size,
181     class_mode='sparse',
182     shuffle=False
183 )
184 test_loss, test_accuracy = model.evaluate(test_ds)
185 print(f"Test Accuracy: {test_accuracy * 100:.2f}%")

```

```

Found 30 images belonging to 6 classes.
1/1 ————— 1s 1s/step - accuracy: 0.5000 - loss: 2.3028
Test Accuracy: 50.00%

```

```

17 [17] model.save("final_model.h5")
18 loaded_model = tf.keras.models.load_model("final_model.h5")

```

WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This file format is considered legacy. We recommend using instead 'model.save\_format="tf"' instead.

WARNING:absl:compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile\_metrics' will be empty until you train or evaluate the model.

```

18 [18] y_true = test_ds.classes
19 y_pred = np.argmax(loaded_model.predict(test_ds), axis=1)
20 print(classification_report(y_true, y_pred, target_names=class_names))

```

```

1/1 ————— 2s 2s/step
precision    recall  f1-score   support

guarana      0.36      1.00      0.53        5
graviola     0.00      0.00      0.00        5
acai          0.40      0.40      0.40        5
cupuacu      0.00      0.00      0.00        5
tucuma       1.00      0.00      0.00        5
pupunha      0.57      0.80      0.67        5

accuracy          0.39      0.50      0.41       30
macro avg         0.39      0.50      0.41       30
weighted avg      0.39      0.50      0.41       30

```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples

\_warn\_prf(average, modifier, f'{metric.capitalize()} is', len(result))

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples

\_warn\_prf(average, modifier, f'{metric.capitalize()} is', len(result))

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples

\_warn\_prf(average, modifier, f'{metric.capitalize()} is', len(result))



```

19 plt.figure(figsize=(12, 6))
20 plt.subplot(1, 2, 1)
21 plt.plot(history.history['accuracy'], label='Train Accuracy')
22 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
23 plt.xlabel('Epochs')
24 plt.ylabel('Accuracy')
25 plt.legend()

26 plt.subplot(1, 2, 2)
27 plt.plot(history.history['loss'], label='Train Loss')
28 plt.plot(history.history['val_loss'], label='Validation Loss')
29 plt.xlabel('Epochs')
30 plt.ylabel('Loss')
31 plt.legend()

32 plt.show()

```



## Task-2

```
base_model = MobileNetV2(input_shape=(img_height, img_width, 3), include_top=False, weights='imagenet')
base_model.trainable = False
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\\_v2/mobilenet\\_v2\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_1.0\\_128\\_no\\_top.h5](https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_128_no_top.h5)  
9406464/9406464 1s 0us/step

```
[21] inputs = Input(shape=(img_height, img_width, 3))
x = base_model(inputs, training=False)
x = GlobalAveragePooling2D()(x)
x = Dense(128, activation='relu')(x)
x = Dropout(0.4)(x)
outputs = Dense(num_classes, activation='softmax')(x)
model = Model(inputs, outputs)

model.summary()
```

Model: "functional\_16"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 128, 128, 3)	0
mobilenetv2_1.00_128 (functional)	(None, 8, 4, 1280)	2,257,280
global_average_pooling2d (GlobalAveragePooling2D)	(None, 1280)	0
dense_2 (Dense)	(None, 128)	163,840
dropout_3 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 0)	128

Total params: 2,421,120 (9.24 MB)  
Trainable params: 163,840 (643.52 KB)  
Non-trainable params: 2,257,280 (8.61 MB)

```
[22] model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```

```
callbacks = [
    ModelCheckpoint("best_model_t1.h5", save_best_only=True, monitor="val_accuracy", mode="max"),
    EarlyStopping(monitor="val_loss", patience=5, restore_best_weights=True),
    ReduceLROnPlateau(monitor="val_loss", factor=0.5, patience=3, min_lr=1e-6)
]

# Train the model (only top layers)
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=30,
    callbacks=callbacks
)
```

Epoch 1/30  
3/3 0s 266ms/step - accuracy: 0.0793 - loss: 3.0859WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This file format is considered legacy. We recommend using instead 'model.save\_model(model)'.  
11s 2s/step - accuracy: 0.0838 - loss: 3.0459 - val\_accuracy: 0.4444 - val\_loss: 1.4776 - learning\_rate: 0.0010

Epoch 2/30  
3/3 0s 598ms/step - accuracy: 0.3889 - loss: 1.6125WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This file format is considered legacy. We recommend using instead 'model.save\_model(model)'.  
3/3 2s 902ms/step - accuracy: 0.3958 - loss: 1.6210 - val\_accuracy: 0.5000 - val\_loss: 1.0801 - learning\_rate: 0.0010

Epoch 3/30  
3/3 0s 593ms/step - accuracy: 0.5852 - loss: 1.1856WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This file format is considered legacy. We recommend using instead 'model.save\_model(model)'.  
3/3 2s 894ms/step - accuracy: 0.5778 - loss: 1.1940 - val\_accuracy: 0.7222 - val\_loss: 0.9430 - learning\_rate: 0.0010

Epoch 4/30  
3/3 2s 734ms/step - accuracy: 0.7000 - loss: 0.9217 - val\_accuracy: 0.6667 - val\_loss: 0.9452 - learning\_rate: 0.0010

Epoch 5/30  
3/3 0s 378ms/step - accuracy: 0.6868 - loss: 0.7543WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This file format is considered legacy. We recommend using instead 'model.save\_model(model)'.  
3/3 2s 684ms/step - accuracy: 0.6818 - loss: 0.7629 - val\_accuracy: 0.8889 - val\_loss: 0.7323 - learning\_rate: 0.0010

Epoch 6/30  
3/3 2s 442ms/step - accuracy: 0.7943 - loss: 0.6844 - val\_accuracy: 0.8889 - val\_loss: 0.5397 - learning\_rate: 0.0010

Epoch 7/30  
3/3 4s 745ms/step - accuracy: 0.8215 - loss: 0.4730 - val\_accuracy: 0.8333 - val\_loss: 0.4606 - learning\_rate: 0.0010

Epoch 8/30  
3/3 2s 439ms/step - accuracy: 0.8394 - loss: 0.3705 - val\_accuracy: 0.8889 - val\_loss: 0.4559 - learning\_rate: 0.0010

Epoch 9/30  
3/3 2s 735ms/step - accuracy: 0.9069 - loss: 0.2700 - val\_accuracy: 0.8889 - val\_loss: 0.5043 - learning\_rate: 0.0010

Epoch 10/30  
3/3 2s 504ms/step - accuracy: 0.8743 - loss: 0.2694 - val\_accuracy: 0.8889 - val\_loss: 0.5666 - learning\_rate: 0.0010

Epoch 11/30  
3/3 2s 689ms/step - accuracy: 0.9510 - loss: 0.1962 - val\_accuracy: 0.8889 - val\_loss: 0.5913 - learning\_rate: 0.0010

Epoch 12/30  
3/3 3s 895ms/step - accuracy: 0.9651 - loss: 0.1586 - val\_accuracy: 0.8889 - val\_loss: 0.5931 - learning\_rate: 5.0000e-04

Epoch 13/30  
3/3 2s 761ms/step - accuracy: 0.9264 - loss: 0.2280 - val\_accuracy: 0.8889 - val\_loss: 0.5772 - learning\_rate: 5.0000e-04

Epoch 14/30  
3/3 2s 761ms/step - accuracy: 0.9264 - loss: 0.2280 - val\_accuracy: 0.8889 - val\_loss: 0.5772 - learning\_rate: 5.0000e-04

[24] test\_loss, test\_accuracy = model.evaluate(test\_ds)  
print(f"Test Accuracy: {test\_accuracy \* 100:.2f}%")

1/1 1s 520ms/step - accuracy: 0.9000 - loss: 0.4587  
Test Accuracy: 90.00%

[25] model.save("final\_model\_t1.h5")  
loaded\_model = tf.keras.models.load\_model("final\_model\_t1.h5")

WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save\_model(model)'. This file format is considered legacy. We recommend using instead 'model.save\_model(model)'.  
WARNING:absl:compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile\_metrics' will be empty until you train or evaluate the model.

✓  
2s

```
▶ y_true = test_ds.classes
  y_pred_probs = loaded_model.predict(test_ds)
  y_pred = np.argmax(y_pred_probs, axis=1)

  print("Inference Output: First 20 Samples:")
  for i in range(20):
      true_label = class_names[int(y_true[i])]
      pred_label = class_names[int(y_pred[i])]
      print(f"{i+1}. True: {true_label} - Predicted: {pred_label}")
```



1/1 2s 2s/step

Inference Output: First 20 Samples:

```
1. True: guarana - Predicted: guarana
2. True: guarana - Predicted: guarana
3. True: guarana - Predicted: acai
4. True: guarana - Predicted: tucuma
5. True: guarana - Predicted: guarana
6. True: graviola - Predicted: graviola
7. True: graviola - Predicted: graviola
8. True: graviola - Predicted: graviola
9. True: graviola - Predicted: graviola
10. True: graviola - Predicted: graviola
11. True: acai - Predicted: acai
12. True: acai - Predicted: acai
13. True: acai - Predicted: acai
14. True: acai - Predicted: acai
15. True: acai - Predicted: acai
16. True: cupuacu - Predicted: cupuacu
17. True: cupuacu - Predicted: cupuacu
18. True: cupuacu - Predicted: cupuacu
```



19. True: cupuacu - Predicted: cupuacu  
20. True: cupuacu - Predicted: cupuacu

✓  
0s

```
[ ] print("Classification Report:")  
    print(classification_report(y_true, y_pred, target_names=class_names))
```

Classification Report:

	precision	recall	f1-score	support
guarana	1.00	0.60	0.75	5
graviola	1.00	1.00	1.00	5
acai	0.83	1.00	0.91	5
cupuacu	0.83	1.00	0.91	5
tucuma	0.83	1.00	0.91	5
pupunha	1.00	0.80	0.89	5
accuracy			0.90	30
macro avg	0.92	0.90	0.89	30
weighted avg	0.92	0.90	0.89	30

```
[28] plt.figure(figsize=(12, 6))  
      plt.subplot(1, 2, 1)  
      plt.plot(history.history['accuracy'], label='Train Accuracy')  
      plt.plot(history.history['val_accuracy'], label='Validation Accuracy')  
      plt.xlabel('Epochs')  
      plt.ylabel('Accuracy')  
      plt.legend()  
  
      plt.subplot(1, 2, 2)  
      plt.plot(history.history['loss'], label='Train Loss')  
      plt.plot(history.history['val_loss'], label='Validation Loss')  
      plt.xlabel('Epochs')  
      plt.ylabel('Loss')  
      plt.legend()  
  
      plt.show()
```

